

How to optimize fuel consumption by using RL?


Overview

- I want an RL agent that learns to control a car (acceleration, deceleration, and gears) and optimize de fuel consumption.
- We have to see the behavior of a car in the real world and try to model it into an RL environment



Modelling

- The agent have to choose between accelerate, don't accelerate, decelerate, up a gear, or down gear.
- Constant acceleration and deceleration
- Real functions to approximate the behavior of a car as much as possible



Revolutions	Gear					
	1	2	3	4	5	6
1000	9,5	16,9	26,3	33,8	43,0	52,6
2000	18,9	33,8	52,6	67,6	86,1	105,2
3000	28,4	50,7	78,9	101,4	129,1	157,8
4000	37,9	67,6	105,2	135,2	172,1	210,3
5000	47,3	84,5	131,5	169,0	215,1	262,9
6000	56,8	101,4	157,8	202,8	258,2	315,5

$$Speed = \frac{C_r \cdot 60 \cdot RPM / 1000}{R_m \cdot RDF}$$

Implementation

- States: Position, speed, gear, revolutions and speed limit
- Actions: Accelerate, don't accelerate, decelerate, gear up and gear down.
- Reward:
$$\begin{cases} -\text{speed} - \text{limit} & \text{if } \text{speed} > \text{limit} \\ -\frac{\text{revolutions}}{\beta_1} - \text{speed} \cdot \beta_2 & \text{if } \text{revolutions} < 2000 \text{ or } \text{revolutions} > 3000 \\ \alpha \cdot \text{speed} - \text{consume} & \text{otherwise} \end{cases}$$
- Update function: **Position:** position += speed*0.2778


Speed: speed += (action - 1) * (6)

Revolutions: revolutions = (self.gear_r[int(gear-1)] * 4.64 * speed * 1000)/(1.83*60)

Obtained from the formula: $\text{Speed} = \frac{C_r \cdot 60 \cdot \text{RPM} / 1000}{R_m \cdot \text{RDF}}$

Consume:


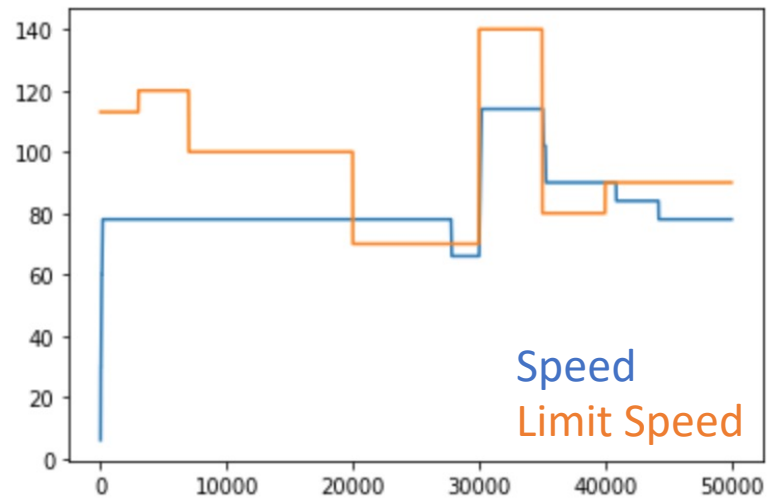
```
if gear == 1:
    consume = 0.0059*speed**2-0.1857*speed+11.0892
elif gear == 2:
    consume = 0.0012*speed**2-0.0574*speed+6.6088
elif gear == 3:
    consume = 0.0005*speed**2-0.0207*speed+4.5582
elif gear == 4:
    consume = 0.0004*speed**2-0.0216*speed+4.1160
elif gear == 5:
    consume = 0.0004*speed**2-0.0183*speed+3.6267
elif gear == 6:
    consume = 0.0004*speed**2-0.0198*speed+3.4347
```



Problems



- Tuning reward function to obtain a consistent result
- Choosing an agent
- Lots of problems with the packages' compatibility
- Very Long execution time

Results



```
: mean_reward_before_train
: {'12345': -44356.8000000000454,
:   '54321': -44356.8000000000454,
:   '12345321': -44356.8000000000454}
```

```
: mean_reward_after_train
: {'12345': 12128.109116994576,
:   '54321': -15145.955977070702,
:   '12345321': -13743.970899999305}
```



Conclusions

- I have managed to create an environment that approximates the behavior of a car in real life
- The results have great variability, hence the importance of reproducibility.
- Next steps: make an environment where you can decide the acceleration that is provided and that is not fixed